

Annand Module 05 Lab 01

Joseph Annand

2023-11-25

Load Libraries

```
library(ISLR2)
library(MASS)
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:ISLR2':
##
## Boston
```

```
library(leaps)
```

```
## Warning: package 'leaps' was built under R version 4.3.2
```

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.2
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(pls)
```

```
## Warning: package 'pls' was built under R version 4.3.2
```

```
##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
## loadings
```

Question 9

Part A - Prepare the Data

```

# Load Dataset
college.data <- College
college.data <- na.omit(college.data)

# Create vector half the size of college.data that contains random set of indices
set.seed(1)
train <- sample(nrow(college.data), 0.8 * nrow(college.data))

# Initialize training and test data
college.train <- college.data[train, ]
college.test <- college.data[-train, ]

```

Part B - Least Squares Regression

```

lm.college <- lm(Apps ~ ., data = college.data, subset = train)

summary(lm.college)

```

```

##
## Call:
## lm(formula = Apps ~ ., data = college.data, subset = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5555.2  -404.6    19.9   310.3  7577.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -630.58238   435.56266  -1.448  0.148209
## PrivateYes  -388.97393   148.87623  -2.613  0.009206 **
## Accept       1.69123    0.04433   38.153 < 2e-16 ***
## Enroll      -1.21543    0.20873  -5.823  9.41e-09 ***
## Top10perc    50.45622    5.88174   8.578 < 2e-16 ***
## Top25perc   -13.62655    4.67321  -2.916  0.003679 **
## F.Undergrad  0.08271    0.03632   2.277  0.023111 *
## P.Undergrad  0.06555    0.03367   1.947  0.052008 .
## Outstate    -0.07562    0.01987  -3.805  0.000156 ***
## Room.Board   0.14161    0.05130   2.760  0.005947 **
## Books        0.21161    0.25184   0.840  0.401102
## Personal     0.01873    0.06604   0.284  0.776803
## PhD          -9.72551    4.91228  -1.980  0.048176 *
## Terminal    -0.48690    5.43302  -0.090  0.928620
## S.F.Ratio    18.26146   13.83984   1.319  0.187508
## perc.alumni  1.39008    4.39572   0.316  0.751934
## Expend       0.05764    0.01254   4.595  5.26e-06 ***
## Grad.Rate    5.89480    3.11185   1.894  0.058662 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 993.8 on 603 degrees of freedom
## Multiple R-squared:  0.9347, Adjusted R-squared:  0.9328

```

```
## F-statistic: 507.5 on 17 and 603 DF,  p-value: < 2.2e-16
```

```
lm.predict <- predict(lm.college, college.test)

lm.mse <- mean((lm.predict - college.test$Apps)^2)
lm.mse
```


```
## [1] 1567324
```

Part C - Ridge Regression

```
# Create matrix of x, the predictors, and vector of y, the response
x <- model.matrix(Apps ~ ., college.data)[, -2]
y <- college.data$Apps

# Create a lambda grid and use it to form ridge regression model
lambda.grid <- 10^seq(10, -2, length = 100)
ridge.mod <- glmnet(x[train, ], y[train], alpha = 0, lambda = lambda.grid, thresh = 1e-12)

# Determine best lambda, or tuning parameter, using cross-validation
set.seed(2)
cv.out <- cv.glmnet(x[train, ], y[train], alpha = 0)
plot(cv.out)
```



```
bestlam <- cv.out$lambda.min
bestlam
```

```
## [1] 362.9786
```


```
# Predict response of test data using ridge regression and calculate MSE
ridge.pred <- predict(ridge.mod, s = bestlam, newx = x[-train, ])
ridge.mse <- mean((ridge.pred - y[-train])^2)
ridge.mse
```

```
## [1] 1508918
```


Part D - The Lasso

```
# Create lasso model on the college dataset
lasso.mod <- glmnet(x[train, ], y[train], alpha = 1, lambda = lambda.grid)
plot(lasso.mod)
```

```
## Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
## collapsing to unique 'x' values
```



```
# Perform cross-validation to determine best lambda or tuning parameter
set.seed(3)
cv.out <- cv.glmnet(x[train, ], y[train], alpha = 1)
plot(cv.out)
```




```
bestlam2<- cv.out$lambda.min

# Predict response of test data and calculate MSE
lasso.pred <- predict(lasso.mod, s = bestlam2, newx = x[-train, ])
lasso.mse <- mean((lasso.pred - y[-train])^2)
lasso.mse
```

```
## [1] 1593402
```

Part E - PCR

```
# Create PCR model on training data
set.seed(4)
pcr.fit <- pcr(Apps ~ ., data = college.data, subset = train, scale = T,
               validation = "CV")
validationplot(pcr.fit, val.type = "MSEP")
axis(side=1, at=seq(1, 20, by=1))
```



Using the cross-validation method, the lowest error occurs when $M=17$.

```
# Predict the number of applications using the PCR model
pcr.pred <- predict(pcr.fit, x[-train, ], ncomp = 17)
pcr.mse <- mean((pcr.pred - y[-train])^2)
pcr.mse
```

```
## [1] 1598902
```


Part F - Partial Least Squares

```
# Create PLS model on the college data
set.seed(5)
pls.fit <- plsr(Apps ~ ., data = college.data, subset = train, scale = T,
  validation = "CV")
summary(pls.fit)
```

```
## Data:      X dimension: 621 17
## Y dimension: 621 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
```

```
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3837    1876    1629    1448    1358    1181    1089
## adjCV           3837    1872    1627    1441    1338    1159    1082
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1073    1066    1060    1059    1056    1056    1055
## adjCV        1068    1062    1056    1054    1051    1051    1050
##      14 comps 15 comps 16 comps 17 comps
## CV          1055    1054    1054    1054
## adjCV        1050    1050    1050    1050
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          25.55   45.38   62.59   65.08   67.55   72.02   75.93   80.46
## Apps       77.30   83.57   87.51   90.88   92.88   93.15   93.24   93.31
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          82.51   85.43   87.83   91.09   92.73   95.12   96.95
## Apps       93.39   93.42   93.45   93.46   93.47   93.47   93.47
##      16 comps 17 comps
## X          97.97   100.00
## Apps       93.47   93.47
```

```
validationplot(pls.fit, val.type = "MSEP")
axis(side=1, at=seq(1, 20, by=1))
```



```
# Predict number of applications using PLS
pls.pred <- predict(pls.fit, x[-train, ], ncomp = 7)
pls.mse <- mean((pls.pred - y[-train])^2)
pls.mse
```

```
## [1] 1528394
```

Part G

```
mse.models <- data.frame(
  model = c("least.squares", "ridge.regression", "lasso", "pcr", "pls"),
  mse = c(lm.mse, ridge.mse, lasso.mse, pcr.mse, pls.mse),
  stringsAsFactors = F
)
mse.models
```

```
##           model      mse
## 1 least.squares 1567324
## 2 ridge.regression 1508918
## 3          lasso 1593402
## 4             pcr 1598902
## 5             pls 1528394
```

Our models adequately explain the relationship between the response and the predictors. the initial least squares regression had an adjusted R-squared of 0.94 and a F-statistic over 500, which indicates that the least squares regression can be used to predict the number of applications received. Several different approaches yield a smaller test error than least squares, meaning that these approaches are even better at predicting the number of applications. The test error for PCR is significantly lower than those of the other four approaches. The other four approaches yield similar test errors to each other.

Question 11

```
boston.data <- Boston
boston.data <- na.omit(boston.data)

# Create vector half the size of college.data that contains random set of indices
set.seed(7)
train <- sample(c(TRUE, FALSE), nrow(boston.data), replace = T)
test <- (!train)
```

Part A - Explore Different Models with Boston data

```
# Best Subset Selection
regfit.full <- regsubsets(crim ~ ., data = boston.data, nvmax = 13)
reg.summary <- summary(regfit.full)
```




```

par(mfrow = c(2,2))
plot(reg.summary$rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
plot(reg.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq",
      type = "l")
which.max(reg.summary$adjr2)

```

```
## [1] 9
```

```
points(9, reg.summary$adjr2[9], col = "red", cex = 2, pch = 20)
```



```
coef(regfit.full, 9)
```

```

##      (Intercept)          zn          indus          nox          dis
## 19.124636156    0.042788127 -0.099385948 -10.466490364 -1.002597606
##          rad          ptratio          black          lstat          medv
## 0.539503547 -0.270835584 -0.008003761 0.117805932 -0.180593877

```

```
# Ridge Regression
```

```

x2 <- model.matrix(crim ~ ., boston.data)[, -1]
y2 <- boston.data$crim

set.seed(8)

```

```
cv.out <- cv.glmnet(x2, y2, alpha = 0)
bestlam3 <- cv.out$lambda.min
```

```
out <- glmnet(x2, y2, alpha = 0)
predict(out, type = "coefficients", s = bestlam3)[1:13,]
```

```
## (Intercept)          zn          indus          chas          nox          rm
## 9.063048666  0.033002416 -0.082046152 -0.737684583 -5.393098508  0.335972073
##          age          dis          rad          tax          ptratio          black
## 0.001962473 -0.702123643  0.422779055  0.003400607 -0.135911587 -0.008483285
##          lstat
## 0.142613436
```

The Lasso


```
set.seed(9)
cv.out <- cv.glmnet(x2, y2, alpha = 1)
bestlam4 <- cv.out$lambda.min

out <- glmnet(x2, y2, alpha = 1, lambda = lambda.grid)
lasso.coef <- predict(out, type = "coefficients", s = bestlam4)[1:13,]
lasso.coef
```

```
## (Intercept)          zn          indus          chas          nox
## 13.3232201680  0.0372265059 -0.0727645312 -0.6003296373 -7.5712246355
##          rm          age          dis          rad          tax
## 0.2670289551  0.0000000000 -0.8227122775  0.5195543611 -0.0001276602
##          ptratio          black          lstat
## -0.2029308099 -0.0075454075  0.1258873002
```

PCR

```
set.seed(11)
pcr.boston <- pcr(crim ~ ., data = boston.data, scale = T,
                  validation = "CV")
validationplot(pcr.boston, val.type = "MSEP")
```



```
# Fit PCR to entire data set using M = 8
pcr.boston <- pcr(y2 ~ x2, scale = T, ncomp = 5)
summary(pcr.boston)
```

```
## Data:      X dimension: 506 13
## Y dimension: 506 1
## Fit method: svdpc
## Number of components considered: 5
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps
## X      47.70   60.36   69.67   76.45   82.99
## y2     30.69   30.87   39.27   39.61   39.61
```

Part B - Propose Model for Predicting Crime Rate

```
# Use Validation-Set Approach to Determine Best Subset Selection Model
regfit.best <- regsubsets(crim ~ ., data = boston.data[train, ], nvmax = 13)

# Create test matrix
test.mat <- model.matrix(crim ~ ., data = boston.data[test, ])

# Compute test MSE for all possible amounts of variables used in the model
val.errors <- rep(NA, 13)
```

```

for (i in 1:13) {
  coefi <- coef(regfit.best, id = i)
  pred <- test.mat[, names(coefi)] %*% coefi
  val.errors[i] <- mean((boston.data$crim[test] - pred)^2)
}

# Get coefficient estimates for model with best subset of variables
best.subset <- which.min(val.errors)
coef(regfit.best, best.subset)

```

```

##      (Intercept)           zn           nox           age           dis
## 21.130421882    0.064159542 -10.950157410    0.030010983   -1.034916464
##           rad           tax           ptratio           black           medv
##  0.705005168   -0.007878952   -0.306715346   -0.005776214   -0.262236529

```

Using the validation set approach on a best subset selection method, a model containing seven predictors was determined to have the lowest MSE of all combinations.

Part C

The chosen model does not involve all features because the best subset selection method was used and with a validation set approach, we determined that the model with the lowest test MSE used only seven of the thirteen possible predictors of crime rate per capita.